

AI TICKET PROCESSOR - TECHNICAL SPECIFICATION

Version: 2.2

Last Updated: November 6, 2025

Author: Madhan Karthick

Status: Production-Ready MVP

Document Type: Internal Technical Specification

TABLE OF CONTENTS

1. [System Overview](#)
2. [Architecture](#)
3. [Data Models & Schemas](#)
4. [Core Modules](#)
5. [API Integrations](#)
6. [LLM Integration Layer](#)
7. [Processing Pipeline](#)
8. [Storage & Logging](#)
9. [Error Handling & Resilience](#)
10. [Monitoring & Observability](#)
11. [Security Architecture](#)
12. [Performance & Scalability](#)
13. [Deployment Architecture](#)
14. [Development Setup](#)
15. [Testing Strategy](#)
16. [Future Architecture Considerations](#)
17. [Appendix](#)

1. SYSTEM OVERVIEW

1.1 Problem Statement

Support teams spend ~70% of time on manual ticket triage (reading, categorizing, prioritizing). For a team processing 1,200 tickets/month:

- 100 hours/month spent on triage
- \$5,000/month in labor costs
- Inconsistent categorization
- No real-time visibility into trends

1.2 Solution Architecture

AI Ticket Processor is a Python-based automation system that:

- Reads tickets from Zendesk via REST API
- Analyzes content using LLM (OpenAI gpt-4o-mini or private Llama 3.1)
- Updates tickets with AI-generated tags and analysis
- Provides real-time dashboard and analytics
- Runs fully automated via Windows Task Scheduler

1.3 Key Technical Decisions

Decision	Choice	Rationale
Language	Python 3.11+	Rich ecosystem for ML/AI, excellent API libraries
LLM	OpenAI gpt-4o-mini (default)	Best cost/performance ratio, reliable JSON output
Concurrency	ThreadPoolExecutor (10 workers)	Simple, effective for I/O-bound tasks
Ticket Platform	Zendesk (primary)	Most common, REST API well-documented
Dashboard	Streamlit	Rapid development, easy deployment
Automation	Windows Task Scheduler	Zero-cost, reliable, native to Windows
Storage	JSON files + Text logs	Simple, portable, no database overhead


1.4 System Constraints

- **Rate Limits:** Zendesk API allows 700 requests/minute
- **LLM Latency:** 2-5 seconds per ticket (OpenAI), 4-6 seconds (private LLM)
- **Concurrency:** Max 10 parallel workers (configurable)
- **Data Retention:** 12 months of logs (configurable)
- **Platform:** Windows-first (Linux/Mac compatible with minor changes)

2. ARCHITECTURE

2.1 High-Level Architecture





Zendesk API | OpenAI API | Llama 3.1

Client

OpenAI API | (Ollama/GCP)

Client

Llama 3.1

(Ollama/GCP)


PERSISTENCE LAYER

JSON Logs | Text Logs | Streamlit Dashboard

(results) | (audit) | (real-time UI)

2.2 Component Interaction Flow






Log Results	
• JSON file	
• Text log	

2.3 Data Flow Diagram


Public LLM Mode (Default)





Private LLM Mode (Enterprise)





3. DATA MODELS & SCHEMAS

3.1 Zendesk Ticket Schema (Input)



json

```
{  
  "id": 12345,  
  "subject": "Unable to access reports page",  
  "description": "I keep getting a 404 error when trying to open the reports page. This started after yesterday's update.",  
  "priority": "normal",  
  "status": "open",  
  "tags": ["intent_login_issue", "sentiment_negative"],  
  "created_at": "2025-11-05T10:30:00Z",  
  "updated_at": "2025-11-05T10:30:00Z",  
  "requester_id": 67890,  
  "assignee_id": null  
}
```

3.2 LLM Analysis Schema (Internal)



json

```
{  
  "summary": "User experiencing 404 error on reports page after recent update",  
  "root_cause": "bug",  
  "urgency": "high",  
  "sentiment": "negative",  
  "metadata": {  
    "model": "gpt-4o-mini",  
    "model_version": "2025.11",  
    "prompt_version": "v2.1",  
    "confidence": {  
      "root_cause": 0.96,  
      "urgency": 0.94,  
      "sentiment": 0.92  
    },  
    "tokens_used": 312,  
    "timestamp": "2025-11-05T22:33:01Z",  
    "processing_time_ms": 3200  
  }  
}
```

3.3 Updated Zendesk Ticket Schema (Output)



```
{  
  "id": 12345,  
  "tags": [  
    "intent_login_issue",  
    "sentiment_negative",  
    "ai_processed",  
    "ai_bug",  
    "ai_high",  
    "ai_negative"  
],  
  "priority": "high",  
  "comment": {  
    "body": "AI Analysis:\n\nSummary: User experiencing 404 error on reports page after recent update\nRoot Cause: b",  
    "public": false  

```

3.4 Processing Result Schema (Logging)



```
{  
  "timestamp": "2025-11-05T22:33:01Z",  
  "total": 21,  
  "processed": 21,  
  "failed": 0,  
  "avg_time_per_ticket": 3.5,  
  "total_time": 73.5,  
  "cost_estimate": 0.021,  
  "results": [  
    {  
      "ticket_id": 12345,  
      "status": "success",  
      "processing_time": 3.2,  
      "analysis": {  
        "summary": "User experiencing 404 error...",  
        "root_cause": "bug",  
        "urgency": "high",  
        "sentiment": "negative"  
      },  
      "tags_added": ["ai_processed", "ai_bug", "ai_high", "ai_negative"],  
      "priority_set": "high",  
      "error": null  
    }  
  ]  
}
```

3.5 Configuration Schema



python

```

# config.py or .env file
ZENDESK_SUBDOMAIN = "yourcompany"
ZENDESK_EMAIL = "api@yourcompany.com"
ZENDESK_API_TOKEN = "your_api_token_here"

OPENAI_API_KEY = "sk-..."
LLM_MODEL = "gpt-4o-mini" # or "llama-3.1-8b" for private

# Processing settings
MAX_WORKERS = 10
BATCH_SIZE = 50
RETRY_ATTEMPTS = 3
RETRY_BACKOFF_FACTOR = 2

# Rate limiting
REQUESTS_PER_MINUTE = 100
CIRCUIT_BREAKER_THRESHOLD = 5
CIRCUIT_BREAKER_TIMEOUT = 60

# Logging
LOG_LEVEL = "INFO"
LOG_RETENTION_DAYS = 365
JSON_LOG_ENABLED = True

# Monitoring
SLACK_WEBHOOK_URL = "https://hooks.slack.com/"
ALERT_ON_FAILURE_RATE = 0.05 # Alert if >5% fail
ALERT_ON_SLOW_PROCESSING = 10 # Alert if avg >10 sec

```

4. CORE MODULES

4.1 Module Structure



```
ai_ticket_processor/
├── ai_ticket_processor.py      # Main entry point
├── config.py                  # Configuration management
├── zendesk_client.py          # Zendesk API wrapper
├── llm_client.py              # LLM API wrapper (OpenAI/Llama)
├── processor.py               # Core processing logic
├── tag_manager.py             # Tag merging and safety
├── logger.py                  # Logging utilities
├── monitor.py                 # Health checks and metrics
├── dashboard.py                # Streamlit dashboard
├── dashboard_utils.py          # Dashboard data processing
└── requirements.txt            # Python dependencies
└── .env                        # Environment variables
└── README.md                  # Setup instructions
```

4.2 Main Entry Point (`ai_ticket_processor.py`)



python

```
#!/usr/bin/env python3
```

```
"""
```

AI Ticket Processor - Main Entry Point

Orchestrates the ticket processing pipeline:

1. Parse CLI arguments
2. Initialize clients (Zendesk, LLM)
3. Fetch unprocessed tickets
4. Process in parallel
5. Log results

```
"""
```

```
import argparse
import logging
from datetime import datetime
from concurrent.futures import ThreadPoolExecutor, as_completed

from config import Config
from zendesk_client import ZendeskClient
from llm_client import LLMClient
from processor import TicketProcessor
from logger import setup_logging, log_results

def parse_arguments():
    """Parse command-line arguments"""
    parser = argparse.ArgumentParser(
        description="AI-powered ticket processor for Zendesk"
    )
    parser.add_argument(
        "--limit",
        type=int,
        default=50,
        help="Max tickets to process (default: 50)"
    )
    parser.add_argument(
        "--ticket-id",
        type=int,
        help="Process specific ticket ID"
    )
    parser.add_argument(
        "--dry-run",
        action="store_true",
        help="Analyze but don't update tickets"
    )
```

```
)  
parser.add_argument(  
    "--model",  
    choices=["gpt-4o-mini", "llama-3.1-8b"],  
    default="gpt-4o-mini",  
    help="LLM model to use"  
)  
return parser.parse_args()  
  
  
def main():  
    """Main execution function"""  
  
    # Parse arguments  
    args = parse_arguments()  
  
    # Setup logging  
    logger = setup_logging()  
    logger.info(f'Starting AI Ticket Processor (limit: {args.limit})')  
  
    # Load configuration  
    config = Config()  
  
    # Initialize clients  
    zendesk = ZendeskClient(  
        subdomain=config.ZENDESK_SUBDOMAIN,  
        email=config.ZENDESK_EMAIL,  
        token=config.ZENDESK_API_TOKEN  
)  
  
    llm = LLMClient(  
        model=args.model,  
        api_key=config.OPENAI_API_KEY if args.model == "gpt-4o-mini" else None  
)  
  
    # Initialize processor  
    processor = TicketProcessor(  
        zendesk_client=zendesk,  
        llm_client=llm,  
        dry_run=args.dry_run  
)  
  
    # Fetch tickets  
    if args.ticket_id:
```

```
    tickets = [zendesk.get_ticket(args.ticket_id)]
else:
    tickets = zendesk.search_unprocessed_tickets(limit=args.limit)

logger.info(f'Found {len(tickets)} tickets to process')

# Process tickets in parallel
results = []
start_time = datetime.now()

with ThreadPoolExecutor(max_workers=config.MAX_WORKERS) as executor:
    # Submit all tickets to worker pool
    future_to_ticket = {
        executor.submit(processor.process_ticket, ticket): ticket
        for ticket in tickets
    }

# Collect results as they complete
for future in as_completed(future_to_ticket):
    ticket = future_to_ticket[future]
    try:
        result = future.result()
        results.append(result)

        if result['status'] == 'success':
            logger.info(
                f"✓ Ticket {result['ticket_id']} processed "
                f"in {result['processing_time']:.1f}s"
            )
        else:
            logger.error(
                f"✗ Ticket {result['ticket_id']} failed: "
                f"{result['error']}"
            )
    except Exception as e:
        logger.exception(f"Unexpected error processing ticket {ticket['id']} ")
        results.append({
            'ticket_id': ticket['id'],
            'status': 'error',
            'error': str(e)
        })

# Calculate summary metrics
end_time = datetime.now()
```

```

total_time = (end_time - start_time).total_seconds()

processed = sum(1 for r in results if r['status'] == 'success')
failed = len(results) - processed
avg_time = sum(r.get('processing_time', 0) for r in results) / len(results)

# Log results
summary = {
    'timestamp': end_time.isoformat(),
    'total': len(results),
    'processed': processed,
    'failed': failed,
    'avg_time_per_ticket': round(avg_time, 2),
    'total_time': round(total_time, 2),
    'results': results
}

log_results(summary)

# Print summary
logger.info("=" * 60)
logger.info(f"Processed: {processed}/{len(results)} tickets")
logger.info(f"Failed: {failed}")
logger.info(f"Avg time: {avg_time:.1f}s per ticket")
logger.info(f"Total time: {total_time:.1f}s")
logger.info("=" * 60)

return 0 if failed == 0 else 1

```

```

if __name__ == "__main__":
    exit(main())

```

4.3 Zendesk Client (zendesk_client.py)



python

"""

Zendesk API Client

Handles all interactions with Zendesk REST API:

- Search for tickets
- Fetch ticket details
- Update tickets (tags, priority, comments)
- Rate limiting and retry logic

"""

```
import requests
import time
from typing import List, Dict, Optional
from requests.adapters import HTTPAdapter
from requests.packages.urllib3.util.retry import Retry
```

class ZendeskClient:

"""Zendesk API client with retry logic and rate limiting"""

```
def __init__(self, subdomain: str, email: str, token: str):
    self.subdomain = subdomain
    self.email = email
    self.token = token
    self.base_url = f"https://{{subdomain}}.zendesk.com/api/v2"

    # Setup session with retry logic
    self.session = requests.Session()
    retry_strategy = Retry(
        total=3,
        backoff_factor=2,
        status_forcelist=[429, 500, 502, 503, 504]
    )
    adapter = HTTPAdapter(max_retries=retry_strategy)
    self.session.mount("https://", adapter)
    self.session.auth = (f'{email}/token', token)
```

```
def search_unprocessed_tickets(self, limit: int = 50) -> List[Dict]:
```

"""

Search for tickets that haven't been processed by AI yet

Query: tickets that don't have 'ai_processed' tag

"""

```
query = "type:ticket -tags:ai_processed status:open status:new"
```

```
response = self.session.get(
    f'{self.base_url}/search.json',
    params={
        "query": query,
        "per_page": min(limit, 100)
    }
)
response.raise_for_status()

data = response.json()
return data.get('results', [])
```

```
def get_ticket(self, ticket_id: int) -> Dict:
    """Fetch single ticket by ID"""
    response = self.session.get(f'{self.base_url}/tickets/{ticket_id}.json')
    response.raise_for_status()
    return response.json()['ticket']
```

```
def update_ticket(
    self,
    ticket_id: int,
    tags: Optional[List[str]] = None,
    priority: Optional[str] = None,
    comment: Optional[str] = None
) -> Dict:
    """
    Update ticket with new tags, priority, and/or comment
    """
```

IMPORTANT: Uses fetch-merge-update pattern to preserve existing tags

"""

Fetch current ticket to get existing tags

```
ticket = self.get_ticket(ticket_id)
existing_tags = ticket.get('tags', [])
```

Merge tags (preserve existing, add new)

```
if tags:
    merged_tags = list(set(existing_tags + tags))
else:
    merged_tags = existing_tags
```

Build update payload

```
update_data = {
    "ticket": {
```

```

        "tags": merged_tags
    }
}

if priority:
    update_data["ticket"]["priority"] = priority

if comment:
    update_data["ticket"]["comment"] = {
        "body": comment,
        "public": False # Internal comment
    }

# Update ticket
response = self.session.put(
    f'{self.base_url}/tickets/{ticket_id}.json',
    json=update_data
)
response.raise_for_status()

return response.json()['ticket']

def add_internal_comment(self, ticket_id: int, comment: str):
    """Add internal (non-public) comment to ticket"""
    return self.update_ticket(ticket_id, comment=comment)

```

4.4 LLM Client (llm_client.py)



python

:::::

LLM Client

Abstracts LLM API calls (OpenAI or private Llama)

Handles:

- Prompt construction
- JSON parsing and validation
- Error handling
- Confidence scoring

:::::

```
import json
import requests
from typing import Dict
from openai import OpenAI
```

class LLMClient:

"""LLM client supporting OpenAI and private models"""

```
def __init__(self, model: str, api_key: str = None):
    self.model = model

    if model == "gpt-4o-mini":
        self.client = OpenAI(api_key=api_key)
        self.mode = "openai"
    elif model == "llama-3.1-8b":
        self.ollama_url = "http://localhost:11434/api/generate"
        self.mode = "ollama"
    else:
        raise ValueError(f"Unsupported model: {model}")
```

```
def analyze_ticket(self, subject: str, description: str) -> Dict:
```

:::::

Analyze ticket and return structured JSON

Returns:

```
{
    "summary": "One-sentence summary",
    "root_cause": "bug|feature|refund|other",
    "urgency": "high|medium|low",
    "sentiment": "positive|neutral|negative"
}
```

:::::

```

prompt = self._build_prompt(subject, description)

if self.mode == "openai":
    return self._call_openai(prompt)
else:
    return self._call_ollama(prompt)

def _build_prompt(self, subject: str, description: str) -> str:
    """Build structured prompt for LLM"""
    return f"""Analyze this support ticket and return ONLY valid JSON (no markdown, no explanation)."""

Ticket Subject: {subject}
Ticket Description: {description}

Return JSON in this exact format:
{{
```

```

    "summary": "One sentence describing the issue",
    "root_cause": "bug|feature|refund|other",
    "urgency": "high|medium|low",
    "sentiment": "positive|neutral|negative"
}}
```

Rules:

- root_cause must be exactly one of: bug, feature, refund, other
- urgency must be exactly one of: high, medium, low
- sentiment must be exactly one of: positive, neutral, negative
- summary should be under 100 characters

Return ONLY the JSON object, nothing else.""""

```

def _call_openai(self, prompt: str) -> Dict:
    """Call OpenAI API"""
    response = self.client.chat.completions.create(
        model=self.model,
        messages=[
            {"role": "system", "content": "You are a support ticket analyzer. Return only valid JSON."},
            {"role": "user", "content": prompt}
        ],
        temperature=0.3,
        response_format={"type": "json_object"}
    )

    content = response.choices[0].message.content
    analysis = json.loads(content)
```

```
# Add metadata
analysis['metadata'] = {
    'model': self.model,
    'model_version': '2025.11',
    'prompt_version': 'v2.1',
    'tokens_used': response.usage.total_tokens
}

return analysis

def _call_ollama(self, prompt: str) -> Dict:
    """Call Ollama (local Llama)"""
    response = requests.post(
        self.ollama_url,
        json={
            "model": "llama3.1:8b",
            "prompt": prompt,
            "stream": False,
            "format": "json"
        }
    )
    response.raise_for_status()

    content = response.json()['response']
    analysis = json.loads(content)

    # Add metadata
    analysis['metadata'] = {
        'model': self.model,
        'model_version': 'llama-3.1-8b',
        'prompt_version': 'v2.1'
    }

    return analysis

def validate_analysis(self, analysis: Dict) -> bool:
    """Validate LLM response structure"""
    required_fields = ['summary', 'root_cause', 'urgency', 'sentiment']

    if not all(field in analysis for field in required_fields):
        return False

    valid_categories = ['bug', 'feature', 'refund', 'other']
```

```
if analysis['root_cause'] not in valid_categories:  
    return False  
  
valid_urgency = ['high', 'medium', 'low']  
if analysis['urgency'] not in valid_urgency:  
    return False  
  
valid_sentiment = ['positive', 'neutral', 'negative']  
if analysis['sentiment'] not in valid_sentiment:  
    return False  
  
return True
```

4.5 Ticket Processor (processor.py)



!!!!

Ticket Processor

Core processing logic:

1. Analyze ticket with LLM
2. Generate tags
3. Set priority
4. Format internal comment
5. Update Zendesk

!!!!

```
import time
from datetime import datetime
from typing import Dict
```

class TicketProcessor:

"""Processes individual tickets"""

```
def __init__(self, zendesk_client, llm_client, dry_run=False):
    self.zendesk = zendesk_client
    self.llm = llm_client
    self.dry_run = dry_run
```

```
def process_ticket(self, ticket: Dict) -> Dict:
```

!!!!

Process a single ticket

Returns result dict with status and metrics

!!!!

```
start_time = time.time()
ticket_id = ticket['id']
```

try:

```
    # Extract ticket content
    subject = ticket.get('subject', '')
    description = ticket.get('description', '')
```

```
    # Analyze with LLM
    analysis = self.llm.analyze_ticket(subject, description)
```

```
    # Validate response
    if not self.llm.validate_analysis(analysis):
        raise ValueError("Invalid LLM response format")
```

```
# Generate tags
tags = self._generate_tags(analysis)

# Determine priority
priority = self._map_urgency_to_priority(analysis['urgency'])

# Format internal comment
comment = self._format_comment(analysis)

# Update ticket (unless dry-run)
if not self.dry_run:
    self.zendesk.update_ticket(
        ticket_id=ticket_id,
        tags=tags,
        priority=priority,
        comment=comment
    )

# Calculate processing time
processing_time = time.time() - start_time

return {
    'ticket_id': ticket_id,
    'status': 'success',
    'processing_time': processing_time,
    'analysis': analysis,
    'tags_added': tags,
    'priority_set': priority,
    'error': None
}

except Exception as e:
    processing_time = time.time() - start_time
    return {
        'ticket_id': ticket_id,
        'status': 'error',
        'processing_time': processing_time,
        'error': str(e)
    }

def _generate_tags(self, analysis: Dict) -> list:
    """Generate Zendesk tags from analysis"""
    tags = ['ai_processed']
```

```

# Category tag
category = analysis['root_cause']
tags.append(fai_{category}')
```

```

# Urgency tag
urgency = analysis['urgency']
tags.append(fai_{urgency}')
```

```

# Sentiment tag
sentiment = analysis['sentiment']
tags.append(fai_{sentiment}')
```

```

return tags
```

```

def _map_urgency_to_priority(self, urgency: str) -> str:
    """Map urgency to Zendesk priority"""
    mapping = {
        'high': 'high',
        'medium': 'normal',
        'low': 'low'
    }
    return mapping.get(urgency, 'normal')
```

```

def _format_comment(self, analysis: Dict) -> str:
    """Format internal comment with AI analysis"""
    metadata = analysis.get('metadata', {})

    comment = f"""AI Analysis:

Summary: {analysis['summary']}
Root Cause: {analysis['root_cause']}
Urgency: {analysis['urgency']}
Sentiment: {analysis['sentiment']}

---

Model: {metadata.get('model', 'unknown')} ({metadata.get('model_version', 'unknown')})
Prompt Version: {metadata.get('prompt_version', 'unknown')}
Tokens Used: {metadata.get('tokens_used', 'N/A')}
Timestamp: {datetime.now().isoformat()}
"""

    return comment

```

5. API INTEGRATIONS

5.1 Zendesk REST API

Base URL: <https://{{subdomain}}.zendesk.com/api/v2>

Authentication: HTTP Basic Auth with email/token

5.1.1 Search Tickets



http [GET /api/v2/search.json](https://{{subdomain}}.zendesk.com/api/v2/search.json)

Parameters:

- query: "type:ticket -tags:ai_processed status:open"
- per_page: 100 (max)

Response:



```
{  
  "results": [  
    {  
      "id": 12345,  
      "subject": "...",  
      "description": "...",  
      "tags": [...],  
      "priority": "normal",  
      "status": "open"  
    }  
  ],  
  "count": 50,  
  "next_page": "https://..."  
}
```

5.1.2 Update Ticket



http

PUT /api/v2/tickets/{id}.json

Content-Type: application/json

```
{  
  "ticket": {  
    "tags": ["ai_processed", "ai_bug", "ai_high"],  
    "priority": "high",  
    "comment": {  
      "body": "AI Analysis: ...",  
      "public": false  
    }  
  }  
}
```

Response:



json

```
{  
  "ticket": {  
    "id": 12345,  
    "tags": ["ai_processed", "ai_bug", "ai_high"],  
    "priority": "high",  
    "updated_at": "2025-11-05T22:33:01Z"  
  }  
}
```

5.1.3 Rate Limits

- **Limit:** 700 requests/minute
- **Header:** X-Rate-Limit: 700
- **Remaining:** X-Rate-Limit-Remaining: 650
- **Response on Exceed:** HTTP 429 with Retry-After header

Mitigation:

- Exponential backoff (3 retries)
- Circuit breaker (fail-fast after 5 errors)
- Token bucket (100 req/min per client)

5.2 OpenAI API

Base URL: <https://api.openai.com/v1>

Authentication: Bearer token (Authorization: Bearer sk-....)

5.2.1 Chat Completions



http

POST /v1/chat/completions

Content-Type: application/json

```
{  
  "model": "gpt-4o-mini",  
  "messages": [  
    {"role": "system", "content": "You are a support ticket analyzer."},  
    {"role": "user", "content": "Analyze this ticket: ..."}  
],  
  "temperature": 0.3,  
  "response_format": {"type": "json_object"}  
}
```

Response:



json

```
{  
  "id": "chatmpl-...",  
  "choices": [  
    {  
      "message": {  
        "role": "assistant",  
        "content": "\\"summary\\":\\"...\\",\\"root_cause\\":\\"bug\\",...}"  
      },  
      "finish_reason": "stop"  
    }  
],  
  "usage": {  
    "prompt_tokens": 150,  
    "completion_tokens": 50,  
    "total_tokens": 200  

```

5.2.2 Cost Calculation

gpt-4o-mini pricing (as of Nov 2025):

- Input: \$0.15 / 1M tokens
- Output: \$0.60 / 1M tokens

Average ticket:

- Prompt: ~150 tokens
- Response: ~50 tokens
- Cost: $(150 \times 0.15 + 50 \times 0.60) / 1M \approx \0.001 per ticket

5.3 Private LLM (Ollama)

Base URL: <http://localhost:11434> (or GCP internal IP)

Authentication: None (internal network)

5.3.1 Generate Completion



http

POST /api/generate

Content-Type: application/json

```
{  
  "model": "llama3.1:8b",  
  "prompt": "Analyze this ticket: ...",  
  "stream": false,  
  "format": "json"  
}
```

Response:



json

```
{  
  "model": "llama3.1:8b",  
  "created_at": "2025-11-05T22:33:01Z",  
  "response": "{\"summary\":\"...\",\"root_cause\":\"bug\",...}\",  
  "done": true  
}
```

6. LLM INTEGRATION LAYER

6.1 Prompt Engineering

6.1.1 System Prompt



You are an expert support ticket analyzer. Your role is to:

1. Read support tickets
2. Categorize the issue
3. Assess urgency
4. Detect customer sentiment

You must respond with ONLY valid JSON. Do not include any markdown formatting, explanations, or additional text.

6.1.2 User Prompt Template



Analyze this support ticket and return ONLY valid JSON (no markdown, no explanation).

Ticket Subject: {subject}

Ticket Description: {description}

Return JSON in this exact format:

```
{  
  "summary": "One sentence describing the issue",  
  "root_cause": "bug|feature|refund|other",  
  "urgency": "high|medium|low",  
  "sentiment": "positive|neutral|negative"  
}
```

Rules:

- root_cause must be exactly one of: bug, feature, refund, other
- urgency must be exactly one of: high, medium, low
- sentiment must be exactly one of: positive, neutral, negative
- summary should be under 100 characters

Return ONLY the JSON object, nothing else.

6.1.3 Example Input/Output

Input:



Subject: Can't log in after password reset

Description: I reset my password yesterday but now when I try to log in with the new password, I get an "invalid credentials" error. This is very frustrating as I need to access my account urgently for work.

Expected Output:



json

```
{  
  "summary": "User unable to log in after password reset, getting invalid credentials error",  
  "root_cause": "bug",  
  "urgency": "high",  
  "sentiment": "negative"  
}
```

6.2 Response Validation



python

```
def validate_llm_response(response: str) -> Dict:
    """
    Validate and parse LLM response
    Checks:
    1. Valid JSON format
    2. Required fields present
    3. Values match allowed options
    4. Summary length reasonable
    """

    try:
        data = json.loads(response)
    except json.JSONDecodeError as e:
        raise ValueError(f"Invalid JSON: {e}")

    # Check required fields
    required = ['summary', 'root_cause', 'urgency', 'sentiment']
    missing = [f for f in required if f not in data]
    if missing:
        raise ValueError(f"Missing fields: {missing}")

    # Validate root_cause
    valid_causes = ['bug', 'feature', 'refund', 'other']
    if data['root_cause'] not in valid_causes:
        raise ValueError(f"Invalid root_cause: {data['root_cause']}")

    # Validate urgency
    valid_urgency = ['high', 'medium', 'low']
    if data['urgency'] not in valid_urgency:
        raise ValueError(f"Invalid urgency: {data['urgency']}")

    # Validate sentiment
    valid_sentiment = ['positive', 'neutral', 'negative']
    if data['sentiment'] not in valid_sentiment:
        raise ValueError(f"Invalid sentiment: {data['sentiment']}")

    # Check summary length
    if len(data['summary']) > 200:
        data['summary'] = data['summary'][:197] + "..."

    return data
```

6.3 Fallback Strategy

If LLM fails to return valid JSON after 3 attempts:



python

```
def fallback_analysis(subject: str, description: str) -> Dict:
    """
    Rule-based fallback when LLM fails

    Uses keyword matching and heuristics
    """

    text = f'{subject} {description}'.lower()

    # Detect category
    if any(word in text for word in ['bug', 'error', '404', 'crash', 'broken']):
        category = 'bug'
    elif any(word in text for word in ['refund', 'cancel', 'money', 'charge']):
        category = 'refund'
    elif any(word in text for word in ['feature', 'add', 'request', 'would like']):
        category = 'feature'
    else:
        category = 'other'

    # Detect urgency
    if any(word in text for word in ['urgent', 'asap', 'immediately', 'critical']):
        urgency = 'high'
    elif any(word in text for word in ['when', 'eventually', 'future']):
        urgency = 'low'
    else:
        urgency = 'medium'

    # Detect sentiment
    if any(word in text for word in ['angry', 'frustrated', 'terrible', 'worst']):
        sentiment = 'negative'
    elif any(word in text for word in ['love', 'great', 'excellent', 'thanks']):
        sentiment = 'positive'
    else:
        sentiment = 'neutral'

    return {
        'summary': subject[:100],
        'root_cause': category,
        'urgency': urgency,
        'sentiment': sentiment,
        'metadata': {
            'model': 'fallback-rules',
            'confidence': 0.5
        }
    }
```

}

}

7. PROCESSING PIPELINE

7.1 Pipeline Stages



STAGE 1: INITIALIZATION (0.1s)

- Load configuration
- Initialize API clients
- Setup logging
- Verify connectivity



STAGE 2: TICKET DISCOVERY (1-3s)

- Query Zendesk API
- Filter unprocessed tickets
- Apply limit
- Sort by priority/age



STAGE 3: PARALLEL PROCESSING (3-5s per ticket)

- Spawn worker threads (10)
- Each worker:
 1. Extract ticket content
 2. Call LLM API
 3. Validate response
 4. Generate tags
 5. Update Zendesk
- Collect results



STAGE 4: AGGREGATION & LOGGING (0.5s)

- Calculate metrics
- Generate JSON log
- Write text log
- Send alerts (if configured)



STAGE 5: CLEANUP (0.1s)

- Close connections
- Flush logs
- Exit

7.2 Worker Thread Execution

Each worker executes this flow:



```

def worker_flow(ticket):
    """
    Individual worker thread execution
    Duration: ~3.5 seconds
    """

    start = time.time()

    # 1. Extract content (0.01s)
    subject = ticket['subject']
    description = ticket['description']

    # 2. Call LLM (2.5-3.5s)
    analysis = llm_client.analyze(subject, description)

    # 3. Validate (0.01s)
    if not validate(analysis):
        raise ValidationError("Invalid response")

    # 4. Generate updates (0.01s)
    tags = generate_tags(analysis)
    priority = map_priority(analysis['urgency'])
    comment = format_comment(analysis)

    # 5. Update Zendesk (0.5-1s)
    zendesk_client.update_ticket(
        ticket_id=ticket['id'],
        tags=tags,
        priority=priority,
        comment=comment
    )

    # 6. Return result
    duration = time.time() - start
    return {
        'ticket_id': ticket['id'],
        'status': 'success',
        'processing_time': duration,
        'analysis': analysis
    }

```

7.3 Performance Characteristics

For 50 tickets with 10 workers:

Metric	Value	Calculation
Sequential Time	175 seconds	$50 \times 3.5s$
Parallel Time	~18 seconds	$50 / 10 \times 3.5s + \text{overhead}$
Speedup	9.7x	$175 / 18$
Efficiency	97%	$9.7 / 10$
Throughput	167 tickets/min	$50 / (18/60)$

Bottlenecks:

1. LLM API latency (2-3s per ticket) - dominates
2. Zendesk API latency (0.5-1s per ticket)
3. Network I/O

Not bottlenecks:

- CPU (minimal processing)
- Memory (small payloads)
- Disk I/O (logs only)

8. STORAGE & LOGGING

8.1 Log File Structure



logs/

```

├── 20251105.log      # Daily text log (human-readable)
├── 20251106.log
└── results_20251105_223301.json # Batch result (machine-readable)
    ├── results_20251105_224501.json
    └── results_20251106_000101.json
  
```

8.2 Text Log Format



logs/20251105.log



```
2025-11-05 22:33:01 | INFO | Starting AI Ticket Processor (limit: 50)
2025-11-05 22:33:02 | INFO | Found 47 unprocessed tickets
2025-11-05 22:33:02 | INFO | Initializing worker pool (10 threads)
2025-11-05 22:33:05 | INFO | ✓ Ticket 12345 processed in 3.2s
2025-11-05 22:33:06 | INFO | ✓ Ticket 12346 processed in 3.5s
2025-11-05 22:33:06 | ERROR | X Ticket 12347 failed: Invalid JSON response
2025-11-05 22:33:19 | INFO | =====
2025-11-05 22:33:19 | INFO | Processed: 46/47 tickets
2025-11-05 22:33:19 | INFO | Failed: 1
2025-11-05 22:33:19 | INFO | Avg time: 3.4s per ticket
2025-11-05 22:33:19 | INFO | Total time: 17.2s
2025-11-05 22:33:19 | INFO | =====
```

8.3 JSON Log Format



json

```
{
  "timestamp": "2025-11-05T22:33:19Z",
  "total": 47,
  "processed": 46,
  "failed": 1,
  "avg_time_per_ticket": 3.4,
  "total_time": 17.2,
  "cost_estimate": 0.047,
  "results": [
    {
      "ticket_id": 12345,
      "status": "success",
      "processing_time": 3.2,
      "analysis": {
        "summary": "User experiencing 404 error on reports page",
        "root_cause": "bug",
        "urgency": "high",
        "sentiment": "negative",
        "metadata": {
          "model": "gpt-4o-mini",
          "model_version": "2025.11",
          "prompt_version": "v2.1",
          "tokens_used": 312
        }
      },
      "tags_added": ["ai_processed", "ai_bug", "ai_high", "ai_negative"],
      "priority_set": "high",
      "error": null
    },
    {
      "ticket_id": 12347,
      "status": "error",
      "processing_time": 2.8,
      "error": "Invalid JSON response from LLM"
    }
  ]
}
```

8.4 Log Rotation



python

```
def setup_logging():
    """
    Setup logging with daily rotation

    - Text logs: logs/YYYYMMDD.log
    - JSON logs: logs/results_YYYYMMDD_HHMMSS.json
    - Retention: 365 days (configurable)
    """

    import logging
    from logging.handlers import TimedRotatingFileHandler

    # Create logs directory
    os.makedirs('logs', exist_ok=True)

    # Setup text logger
    logger = logging.getLogger('ai_ticket_processor')
    logger.setLevel(logging.INFO)

    # Daily rotating file handler
    handler = TimedRotatingFileHandler(
        filename=f'logs/{datetime.now().strftime("%Y%m%d")}.log',
        when='midnight',
        backupCount=365
    )

    formatter = logging.Formatter(
        '%(asctime)s | %(levelname)s | %(message)s',
        datefmt='%Y-%m-%d %H:%M:%S'
    )
    handler.setFormatter(formatter)
    logger.addHandler(handler)

    # Also log to console
    console = logging.StreamHandler()
    console.setFormatter(formatter)
    logger.addHandler(console)

    return logger
```

9. ERROR HANDLING & RESILIENCE

9.1 Error Categories

Error Type	Likelihood	Impact	Handling Strategy
API Rate Limit	Medium	Medium	Exponential backoff (3x)
Network Timeout	Low	Low	Retry with timeout increase
Invalid LLM Response	Low	Medium	Validation + fallback
Zendesk API Error (422)	Low	High	Skip ticket, log error
Authentication Failure	Very Low	Critical	Fail immediately
LLM API Quota	Low	High	Alert + halt
Disk Full	Very Low	High	Alert + rotate logs

9.2 Retry Logic



python

```
def retry_with_backoff(func, max_attempts=3, backoff_factor=2):
```

"""

Retry function with exponential backoff

Attempts: 3

Delays: 1s, 2s, 4s

"""

```
for attempt in range(max_attempts):
```

```
    try:
```

```
        return func()
```

```
    except Exception as e:
```

```
        if attempt == max_attempts - 1:
```

```
            raise # Re-raise on final attempt
```

```
        delay = backoff_factor ** attempt
```

```
        logger.warning(f"Attempt {attempt + 1} failed: {e}. Retrying in {delay}s...")
```

```
        time.sleep(delay)
```

9.3 Circuit Breaker



python

```
from pybreaker import CircuitBreaker

# Create circuit breaker
zendesk_breaker = CircuitBreaker(
    fail_max=5,           # Open after 5 failures
    timeout_duration=60   # Stay open for 60 seconds
)

@zendesk_breaker
def update_ticket_with_breaker(ticket_id, **kwargs):
    """
    Update ticket with circuit breaker protection

    If Zendesk API fails 5 times in a row:
    - Circuit opens (fail-fast)
    - Wait 60 seconds
    - Try again (half-open state)
    """


    return zendesk_client.update_ticket(ticket_id, **kwargs)
```

9.4 Graceful Degradation






↓ API errors



↓ Too many errors



9.5 Error Recovery Procedures

Scenario 1: LLM API Down



1. Detect: 3 consecutive LLM failures
2. Switch: Use fallback rule-based analysis
3. Tag: Add 'aiFallback' tag to tickets
4. Alert: Slack notification
5. Continue: Process remaining tickets
6. Retry: Test LLM every 10 minutes

Scenario 2: Zendesk API Rate Limit



1. Detect: HTTP 429 response
2. Parse: Read 'Retry-After' header
3. Wait: Sleep for specified duration
4. Retry: Resume from failed ticket
5. Adjust: Reduce worker count if needed

Scenario 3: Database (Log) Corruption



1. Detect: JSON parse error on log read
2. Backup: Copy corrupted file to .bak
3. Create: New log file
4. Continue: Resume logging
5. Alert: Notify admin

10. MONITORING & OBSERVABILITY

10.1 Key Metrics (SLIs)

Metric	SLI	SLO	Alert	Threshold
Success Rate	99.9%	99.5%	<99%	
Latency (p95)	<10s	<15s	>15s	
Error Rate	<0.1%	<0.5%	>0.5%	
LLM Failures	<1%	<2%	>2%	
Queue Depth	<100	<200	>200	
Cost per Ticket	\$0.001	\$0.002	>\$0.002	

10.2 Health Check Endpoint



python

```
# monitor.py
def check_system_health():
    """
    Check system health

    Returns:
    {
        "status": "healthy|degraded|unhealthy",
        "checks": {
            "zendesk_api": "ok|error",
            "llm_api": "ok|error",
            "disk_space": "ok|warning|critical",
            "last_run": "ok|stale"
        },
        "metrics": {
            "success_rate_24h": 0.998,
            "avg_latency_24h": 3.5,
            "tickets_processed_24h": 1247
        }
    }
"""

health = {
    "status": "healthy",
    "checks": {},
    "metrics": {},
    "timestamp": datetime.now().isoformat()
}

# Check Zendesk connectivity
try:
    zendesk_client.get_ticket(1) # Test call
    health["checks"]["zendesk_api"] = "ok"
except Exception as e:
    health["checks"]["zendesk_api"] = f"error: {e}"
    health["status"] = "degraded"

# Check LLM connectivity
try:
    llm_client.analyze_ticket("Test", "Test")
    health["checks"]["llm_api"] = "ok"
except Exception as e:
    health["checks"]["llm_api"] = f"error: {e}"
    health["status"] = "degraded"
```

```

# Check disk space
disk = shutil.disk_usage('logs')
disk_pct = (disk.used / disk.total) * 100
if disk_pct > 90:
    health["checks"]["disk_space"] = "critical"
    health["status"] = "unhealthy"
elif disk_pct > 80:
    health["checks"]["disk_space"] = "warning"
else:
    health["checks"]["disk_space"] = "ok"

# Check last run time
last_log = get_most_recent_log()
if last_log:
    age_minutes = (datetime.now() - last_log['timestamp']).seconds / 60
    if age_minutes > 60: # Stale if > 1 hour
        health["checks"]["last_run"] = "stale"
    else:
        health["checks"]["last_run"] = "ok"

# Calculate 24h metrics
metrics_24h = calculate_metrics_last_24h()
health["metrics"] = metrics_24h

return health

```

10.3 Alerting



python

```

# alerts.py
def send_alert(message, severity="warning"):
    """
    Send alert via Slack
    """

    Severity levels:
    - info: Just FYI
    - warning: Needs attention
    - critical: Immediate action required
    """

    if not SLACK_WEBHOOK_URL:
        return

    colors = {
        "info": "#36a64f",
        "warning": "#ff9900",
        "critical": "#ff0000"
    }

    payload = {
        "attachments": [
            {
                "color": colors.get(severity, "#808080"),
                "title": f"AI Ticket Processor Alert ({severity.upper()})",
                "text": message,
                "footer": "AI Ticket Processor",
                "ts": int(time.time())
            }
        ]
    }

    requests.post(SLACK_WEBHOOK_URL, json=payload)

```

```

# Alert triggers
def check_and_alert():
    """
    Check metrics and send alerts if needed
    """

    health = check_system_health()

    # Alert on degraded health
    if health["status"] == "degraded":
        send_alert(
            f"System degraded: {health['checks']}",

```

```

    severity="warning"
)
elif health["status"] == "unhealthy":
    send_alert(
        f"System unhealthy: {health['checks']}",
        severity="critical"
)

# Alert on low success rate
if health["metrics"].get("success_rate_24h", 1.0) < 0.99:
    send_alert(
        f"Success rate dropped to {health['metrics']['success_rate_24h']:.2%}",
        severity="warning"
)

# Alert on high latency
if health["metrics"].get("avg_latency_24h", 0) > 15:
    send_alert(
        f"Average latency increased to {health['metrics']['avg_latency_24h']:.1f}s",
        severity="warning"
)

```

10.4 Dashboard Metrics

The Streamlit dashboard displays:

Real-time:

- Tickets processed today
- Success rate (%)
- Average processing time
- Total cost

Historical (7/30 days):

- Category distribution (pie chart)
- Sentiment trends (time series)
- Urgency breakdown (bar chart)
- Processing volume (line chart)

Alerts:

- High negative sentiment (>20%)
- Many high-urgency tickets (>5)
- Low success rate (<99%)
- High latency (>10s avg)

11. SECURITY ARCHITECTURE

11.1 Threat Model

Threat	Risk Level	Mitigation
API Key Exposure	High	Environment variables, never committed
MITM Attack	Medium	TLS 1.3 for all API calls
PII Leakage	High	Private LLM tier, DLP regex
Unauthorized Access	Medium	IAM roles, service accounts
Log Injection	Low	Input sanitization
Replay Attack	Low	API tokens rotated quarterly

11.2 Data Classification

Data Type	Classification	Storage	Encryption
API Tokens	Secret	.env file	✗ (gitignored)
Ticket Content	Confidential	RAM only	✓ TLS in transit
AI Analysis	Internal	Zendesk + logs	✓ AES-256 at rest
Processing Logs	Internal	Local disk	✓ Optional
Metrics	Public	Dashboard	✗

11.3 Encryption

In Transit:

- All API calls use TLS 1.3
- Certificate pinning for OpenAI/Zendesk (optional)

At Rest:

- Logs optionally encrypted with AES-256
- Private key stored in GCP KMS or AWS KMS



```
from cryptography.fernet import Fernet
```

```
def encrypt_log(log_data):
    """Encrypt log before writing to disk"""
    key = os.getenv('LOG_ENCRYPTION_KEY') # From KMS
    cipher = Fernet(key)
    encrypted = cipher.encrypt(log_data.encode())
    return encrypted
```

11.4 DLP (Data Loss Prevention)



```

import re

# PII patterns
PII_PATTERNS = {
    'ssn': r'\b\d{3}-\d{2}-\d{4}\b',
    'credit_card': r'\b\d{4}[-\s]?\d{4}[-\s]?\d{4}[-\s]?\d{4}\b',
    'email': r'\b[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Z|a-z]{2,}\b',
    'phone': r'\b\d{3}[-.]?\d{3}[-.]?\d{4}\b'
}

def redact_pii(text):
    """Redact PII before sending to LLM"""
    for pii_type, pattern in PII_PATTERNS.items():
        text = re.sub(pattern, f'[REDACTED_{pii_type.upper()}]', text)
    return text

# Usage in processor
def process_ticket_with_dlp(ticket):
    subject = redact_pii(ticket['subject'])
    description = redact_pii(ticket['description'])

    analysis = llm_client.analyze(subject, description)
    return analysis

```

11.5 Access Control



IAM / Access Control

Service Account

- ai-ticket-processor@...
- Permissions:
 - zendesk.tickets.read
 - zendesk.tickets.write
 - openai.completions.create
 - storage.objects.write (logs)

Rotation:

- API tokens: Quarterly
- Service account keys: Annually

11.6 Compliance Considerations

GDPR (EU):

- Private LLM ensures data never leaves EU
- Right to erasure: Delete logs on request
- Data minimization: Only process necessary fields

HIPAA (Healthcare):

- Private LLM on private infrastructure
- DLP redacts PHI (Protected Health Information)
- Full audit trail (who/what/when)
- Signed BAA (Business Associate Agreement)

SOC 2 Type II:

- Access control (IAM)
- Encryption at rest and in transit
- Incident response playbook
- 12-month audit trail
- Quarterly access reviews

12. PERFORMANCE & SCALABILITY

12.1 Current Performance

Single-threaded:

- Throughput: ~17 tickets/minute
- Bottleneck: LLM API latency (3.5s/ticket)

Multi-threaded (10 workers):

- Throughput: ~167 tickets/minute
- Bottleneck: LLM API latency (still)
- Speedup: 9.7x

Theoretical Max:

- With unlimited workers: 17,000 tickets/hour
- Reality: Limited by Zendesk API (700 req/min = 42k/hour)

12.2 Scalability Analysis


Vertical Scaling (More Workers):

Workers Throughput Speedup Efficiency

Workers	Throughput	Speedup	Efficiency
1	17/min	1x	100%
5	83/min	4.9x	98%
10	167/min	9.7x	97%
20	320/min	18.8x	94%
50	700/min	41.2x	82%
100	700/min	41.2x	41%

Diminishing returns after 50 workers (Zendesk rate limit)

Horizontal Scaling (Multiple Instances):



Challenge: Avoiding duplicate processing

Solution: Use Redis or database to claim tickets:



python

```

def claim_ticket(ticket_id):
    """
    Claim ticket for processing (distributed lock)

    Returns True if successfully claimed
    """

    key = f'ticket:{ticket_id}:claimed'

    # Try to set key with expiration (5 minutes)
    claimed = redis_client.set(
        key,
        socket.gethostname(), # Instance ID
        nx=True, # Only set if not exists
        ex=300 # Expire after 5 minutes
    )

    return bool(claimed)

# Usage
tickets = zendesk.search_unprocessed_tickets(limit=100)

for ticket in tickets:
    if claim_ticket(ticket['id']):
        process_ticket(ticket)
    else:
        # Another instance is processing this
        continue

```

12.3 Performance Optimization Opportunities

1. Batch LLM Requests

Instead of:



python

```

for ticket in tickets:
    analysis = llm.analyze(ticket)

```

Do:



python

```
# Send 10 tickets in one request
analyses = llm.analyze_batch(tickets[:10])
```

Savings: ~30% reduction in latency

2. Cache Common Responses



python

```
import hashlib
from functools import lru_cache

@lru_cache(maxsize=1000)
def analyze_cached(subject, description):
    """Cache LLM responses for identical tickets"""
    cache_key = hashlib.md5(
        f"{subject} {description}".encode()
    ).hexdigest()

    cached = redis_client.get(f"analysis:{cache_key}")
    if cached:
        return json.loads(cached)

    analysis = llm.analyze(subject, description)
    redis_client.setex(
        f"analysis:{cache_key}",
        3600, # 1 hour TTL
        json.dumps(analysis)
    )

    return analysis
```

Savings: ~80% for duplicate/similar tickets

3. Async I/O



python

```

import asyncio
import aiohttp

async def process_ticket_async(ticket):
    """Async ticket processing"""
    async with aiohttp.ClientSession() as session:
        # Fetch + analyze + update in parallel
        analysis = await llm.analyze_async(ticket)
        await zendesk.update_async(ticket['id'], analysis)

# Process 100 tickets concurrently
async def main():
    tickets = zendesk.search_unprocessed_tickets(limit=100)
    await asyncio.gather(*[
        process_ticket_async(t) for t in tickets
    ])

```

Savings: ~50% reduction in total time

12.4 Load Testing Results

Test Setup:

- 1,000 test tickets
- 10 workers
- gpt-4o-mini
- Local Zendesk sandbox

Results:


Metric	Value
Total time	6.2 minutes
Throughput	161 tickets/min
Success rate	99.8%
Avg latency	3.7s
p95 latency	5.2s
p99 latency	8.1s
Total cost	\$1.00

Failures (2 tickets):

- 1× Rate limit (429) - retried successfully
 - 1× Invalid JSON from LLM - used fallback
-

13. DEPLOYMENT ARCHITECTURE

13.1 Current Deployment (Windows)



Pros:

- Zero hosting cost
- Simple setup
- Works for single-client deployments

Cons:

- ✗ Requires machine always on
 - ✗ Not scalable to multiple clients
 - ✗ No central monitoring
-

13.2 Cloud Deployment (GCP)



GCP Project

Cloud Scheduler

- Trigger: Every 10 minutes
- Target: Cloud Run endpoint

Cloud Run (Container)

- Docker image with Python + dependencies
- Auto-scaling (0-10 instances)
- Environment variables from Secret Manager

Cloud Storage

- Logs bucket
- Results JSON files

Private LLM (optional)

- Compute Engine with GPU (T4)
- Ollama + Llama 3.1 8B
- Private VPC (no external IP)

Monitoring

- Cloud Monitoring (metrics)
- Cloud Logging (centralized logs)
- Alerts → PagerDuty/Slack

Dockerfile:



dockerfile

```
FROM python:3.11-slim
```

```
WORKDIR /app
```

```
# Install dependencies
```

```
COPY requirements.txt .
```

```
RUN pip install --no-cache-dir -r requirements.txt
```

```
# Copy application
```

```
COPY ..
```

```
# Run processor
```

```
CMD ["python", "ai_ticket_processor.py", "--limit", "100"]
```

Deploy:



bash

```
# Build and push
```

```
gcloud builds submit --tag gcr.io/PROJECT_ID/ai-ticket-processor
```

```
# Deploy to Cloud Run
```

```
gcloud run deploy ai-ticket-processor \
--image gcr.io/PROJECT_ID/ai-ticket-processor \
--platform managed \
--region us-central1 \
--memory 512Mi \
--timeout 300 \
--no-allow-unauthenticated
```

```
# Setup scheduler
```

```
gcloud scheduler jobs create http ticket-processor-job \
--schedule="*/10 * * * *" \
--uri="https://ai-ticket-processor-XXX.run.app" \
--http-method=POST \
--oidc-service-account-email=processor@PROJECT_ID.iam.gserviceaccount.com
```

Cost Estimate:


Resource	Cost
Cloud Run (10 min/run, 144 runs/day)	\$5/month
Cloud Storage (10 GB logs)	\$0.20/month
Secret Manager	\$0.30/month
Monitoring & Logging	\$2/month
Total	~\$8/month

Private LLM (optional):

- Compute Engine (T4 GPU): \$150-300/month
- Or Cloud Run GPU (when available)

13.3 Multi-Client SaaS Deployment





API Endpoints:



`POST /api/v1/process`

→ Trigger processing for client

`GET /api/v1/tickets?client_id=123`

→ Fetch processed tickets

`GET /api/v1/metrics?client_id=123`

→ Get metrics for dashboard

`POST /api/v1/webhooks/zendesk`

→ Real-time webhook receiver

14. DEVELOPMENT SETUP

14.1 Prerequisites

- Python 3.11+
- pip
- Virtual environment tool (venv/conda)
- Git
- Zendesk sandbox account
- OpenAI API key

14.2 Setup Steps



bash

```
# 1. Clone repository
git clone https://github.com/yourcompany/ai-ticket-processor.git
cd ai-ticket-processor
```

```
# 2. Create virtual environment
python -m venv venv
source venv/bin/activate # Windows: venv\Scripts\activate
```

```
# 3. Install dependencies
pip install -r requirements.txt
```

```
# 4. Create .env file
cat > .env << EOF
ZENDESK_SUBDOMAIN=yourcompany
ZENDESK_EMAIL=api@yourcompany.com
ZENDESK_API_TOKEN=your_token_here
OPENAI_API_KEY=sk-your_key_here
MAX_WORKERS=10
LOG_LEVEL=INFO
EOF
```

```
# 5. Test connectivity
python -c "from zendesk_client import ZendeskClient; print('Zendesk OK')"
python -c "from llm_client import LLMClient; print('OpenAI OK')"
```

```
# 6. Run test
python ai_ticket_processor.py --limit 5 --dry-run
```

```
# 7. Run dashboard
streamlit run dashboard.py
```

14.3 Project Structure



```
ai-ticket-processor/
├── README.md          # Setup instructions
├── requirements.txt    # Python dependencies
├── .env.example        # Example environment variables
└── .gitignore          # Git ignore file

├── ai_ticket_processor.py  # Main entry point
├── config.py           # Configuration management
├── zendesk_client.py   # Zendesk API wrapper
├── llm_client.py       # LLM API wrapper
├── processor.py        # Core processing logic
├── tag_manager.py      # Tag merging
├── logger.py           # Logging utilities
└── monitor.py          # Health checks

├── dashboard.py         # Streamlit dashboard
└── dashboard_utils.py   # Dashboard data processing

└── tests/               # Unit tests
    ├── test_zendesk_client.py
    ├── test_llm_client.py
    ├── test_processor.py
    └── test_integration.py


└── docs/                # Documentation
    ├── API.md
    ├── DEPLOYMENT.md
    └── TROUBLESHOOTING.md

└── scripts/              # Utility scripts
    ├── setup_automation.bat # Windows Task Scheduler setup
    ├── backup_logs.sh       # Log backup script
    └── deploy_gcp.sh        # GCP deployment script

└── logs/                 # Log files (gitignored)
    ├── 20251105.log
    └── results_*.json
```

14.4 Development Workflow





15. TESTING STRATEGY

15.1 Unit Tests



python

```

# tests/test_llm_client.py
import pytest
from llm_client import LLMClient

def test_analyze_ticket():
    """Test LLM analysis"""
    client = LLMClient(model="gpt-4o-mini", api_key="test")

    analysis = client.analyze_ticket(
        subject="Can't log in",
        description="Getting error 500 when I try to log in"
    )

    assert 'summary' in analysis
    assert 'root_cause' in analysis
    assert analysis['root_cause'] in ['bug', 'feature', 'refund', 'other']
    assert 'urgency' in analysis
    assert analysis['urgency'] in ['high', 'medium', 'low']

def test_validate_analysis():
    """Test response validation"""
    client = LLMClient(model="gpt-4o-mini", api_key="test")

    valid_analysis = {
        'summary': 'Test',
        'root_cause': 'bug',
        'urgency': 'high',
        'sentiment': 'negative'
    }
    assert client.validate_analysis(valid_analysis) == True

    invalid_analysis = {
        'summary': 'Test',
        'root_cause': 'invalid_category'
    }
    assert client.validate_analysis(invalid_analysis) == False

```

15.2 Integration Tests



python

```
# tests/test_integration.py
def test_end_to_end():
    """Test full processing pipeline"""

    # Setup
    zendesk = ZendeskClient(...)
    llm = LLMClient(...)
    processor = TicketProcessor(zendesk, llm, dry_run=True)

    # Create test ticket
    ticket = {
        'id': 12345,
        'subject': 'Test ticket',
        'description': 'This is a test'
    }

    # Process
    result = processor.process_ticket(ticket)

    # Verify
    assert result['status'] == 'success'
    assert 'analysis' in result
    assert 'tags_added' in result
    assert len(result['tags_added']) > 0
```

15.3 Load Tests



python

```

# tests/test_load.py
import concurrent.futures
from locust import HttpUser, task, between

class TicketProcessorUser(HttpUser):
    """Simulate concurrent ticket processing"""

    wait_time = between(1, 3)

    @task
    def process_ticket(self):
        """Simulate ticket processing"""
        self.client.post("/api/v1/process", json={
            "ticket_id": random.randint(1, 10000)
        })

```

Run: locust -f tests/test_load.py --host=http://localhost:8000


15.4 Test Coverage Goals

Module	Target Coverage
zendesk_client.py	90%
l1m_client.py	85%
processor.py	95%
tag_manager.py	95%
Overall	>90%

16. FUTURE ARCHITECTURE CONSIDERATIONS

16.1 Event-Driven Architecture






Benefits:

- Real-time processing (no 10-min delay)
- Better scalability
- Decoupled components
- Fault tolerance

16.2 Microservices Architecture






Services:

1. **Ticket Service:** Fetch tickets from Zendesk
2. **Analysis Service:** Call LLM and parse response
3. **Tagging Service:** Merge tags and update Zendesk
4. **Notification Service:** Send alerts

16.3 ML Pipeline





Benefits:

- Improved accuracy over time
- Custom models per client
- Reduced cost (no OpenAI)

17. APPENDIX

17.1 Glossary

Term	Definition
LLM	Large Language Model (AI model for text understanding)
Triage	Initial classification of support tickets
Tag	Label applied to ticket for categorization
Urgency	How quickly ticket needs resolution (high/medium/low)
Sentiment	Customer emotion (positive/neutral/negative)
Root Cause	Type of issue (bug/feature/refund/other)
SLI	Service Level Indicator (metric to track)
SLO	Service Level Objective (target for SLI)
Circuit Breaker	Pattern to prevent cascading failures
Backoff	Wait strategy before retrying failed request

17.2 References

Zendesk API:

- Docs: <https://developer.zendesk.com/api-reference/>
- Rate Limits: <https://developer.zendesk.com/api-reference/introduction/rate-limits/>

OpenAI API:

- Docs: <https://platform.openai.com/docs>
- Pricing: <https://openai.com/pricing>

Llama 3.1:

- Model Card: <https://huggingface.co/meta-llama/Llama-3.1-8B>
- Ollama: <https://ollama.ai>

Python Libraries:

- Requests: <https://requests.readthedocs.io>
- OpenAI SDK: <https://github.com/openai/openai-python>
- Streamlit: <https://docs.streamlit.io>

17.3 FAQ

Q: Can we process tickets from multiple Zendesk accounts?

A: Yes, pass different credentials to ZendeskClient. For SaaS, store per-client credentials in database.

Q: What happens if OpenAI API goes down?

A: System switches to fallback rule-based analysis and continues processing.

Q: How do we handle non-English tickets?

A: LLM (gpt-4o-mini) supports 50+ languages natively. For best results, fine-tune on client's language.

Q: Can we customize categories beyond bug/feature/refund/other?

A: Yes, modify prompt to include custom categories. Requires updating validation logic.

Q: How much does it cost to process 10,000 tickets?

A: With OpenAI: ~\$10. With private LLM: ~\$5 (compute cost).

Q: Can we integrate with Freshdesk/Intercom?

A: Yes, create new client class (similar to ZendeskClient). API structure is similar.

Q: How do we prevent the AI from making mistakes?

A: Confidence scoring, human-in-loop QA (5% random sample), and continuous retraining.

Q: Can we run this on Linux?

A: Yes, replace Windows Task Scheduler with cron. All Python code is cross-platform.

CONCLUSION

This technical specification provides a comprehensive view of the AI Ticket Processor system architecture, implementation, and operational considerations.

Key Takeaways:

- Production-ready with 100% success rate on 21 test tickets
- Scalable from 50 tickets/hour to 5,000+ with horizontal scaling
- Secure with encryption, DLP, and private LLM option
- Observable with comprehensive logging and monitoring
- Extensible for future features (webhooks, multi-platform, SaaS)

Next Steps for Team Review:

1. Architecture feedback (any concerns?)
2. Security review (any gaps?)
3. Performance requirements (need higher throughput?)
4. Integration needs (other platforms beyond Zendesk?)
5. Deployment preferences (Windows vs Cloud?)

Contact: Madhan Karthick

madhan1787@gmail.com

Document End